

Process Environment

Advanced Programming Environment in the UNIX Environment

Chun-Ying Huang <chuang@cs.nctu.edu.tw>

Outline

Process start and termination

Environment variables

Memory layout

Shared libraries

Memory allocation

setjmp and longjmp

Process resource limits

Process Start

The main function

Synopsis

- `int main(int argc, char *argv[]);`
- `int main(int argc, char *argv[], char *envp[]);`

Process Termination

Normal process termination in five ways

- Return from main
- Calling exit
- Calling _exit or _Exit
- Return of the last thread from its start routine
- Calling pthread_exit from the last thread

Abnormal process termination in three ways

- Calling abort
- Receipt of a signal
- Response of the last thread to a cancellation request

Execution of a main function looks like

- `exit(main(argc, argv));`

atexit and exit Functions

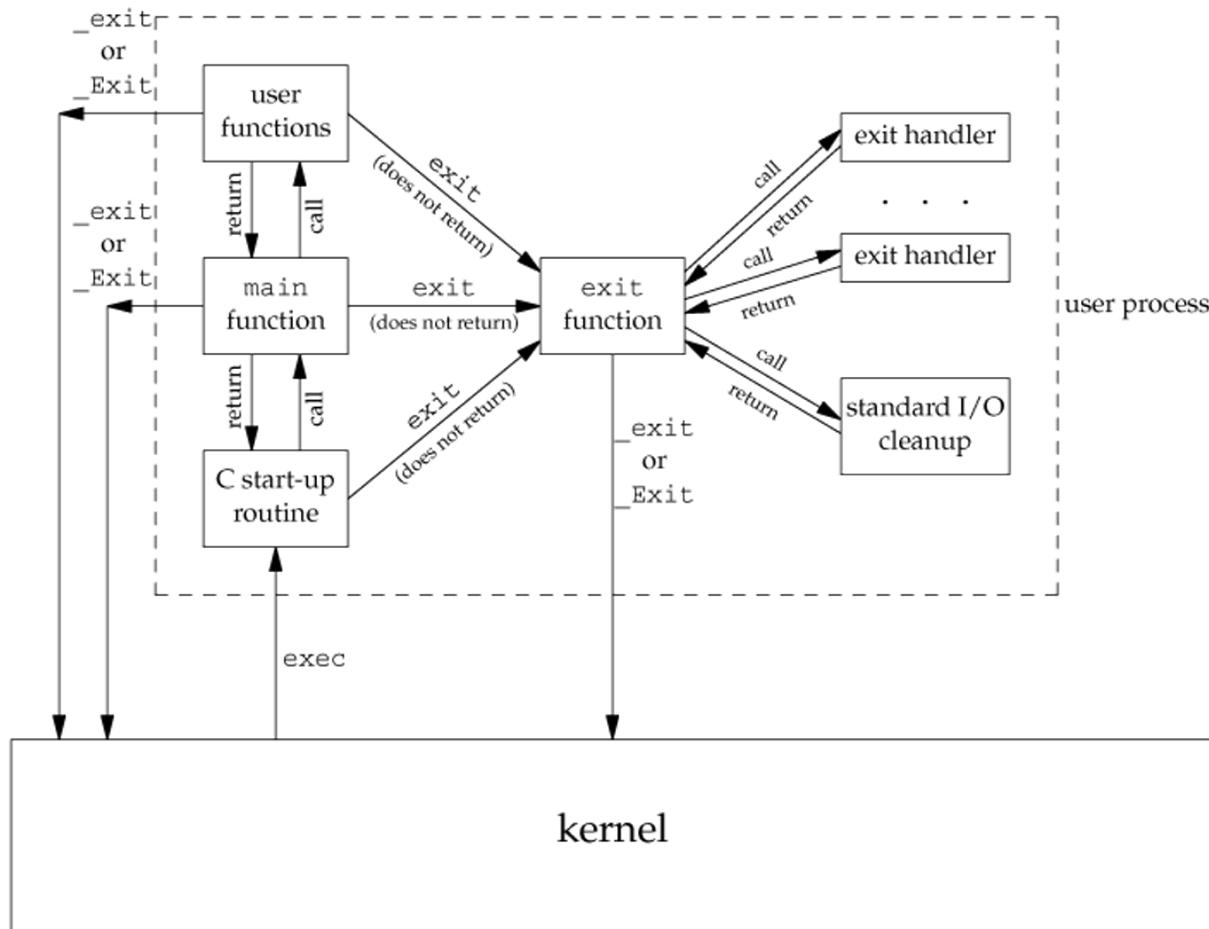
Manual cleanups on exit

- `int atexit(void (*function)(void));`
- Register up to 32 customized functions (textbook)
 - Linux has extended this restrictions

Exit functions

- `exit`
 - Call atexit registered functions
 - Performed a clean shutdown of the standard I/O library
 - `fclose()` all streams, `remove tmpfile()`
- `_exit` and `_Exit`
 - Terminate immediately

Start and Termination of a C Program



Environment Variables

The environment variables

- Usually in the form of: **name=value** (no spaces around =)
- Relevant commands: env, export (bash)
- Use \$ to read a specific environment variable in a shell

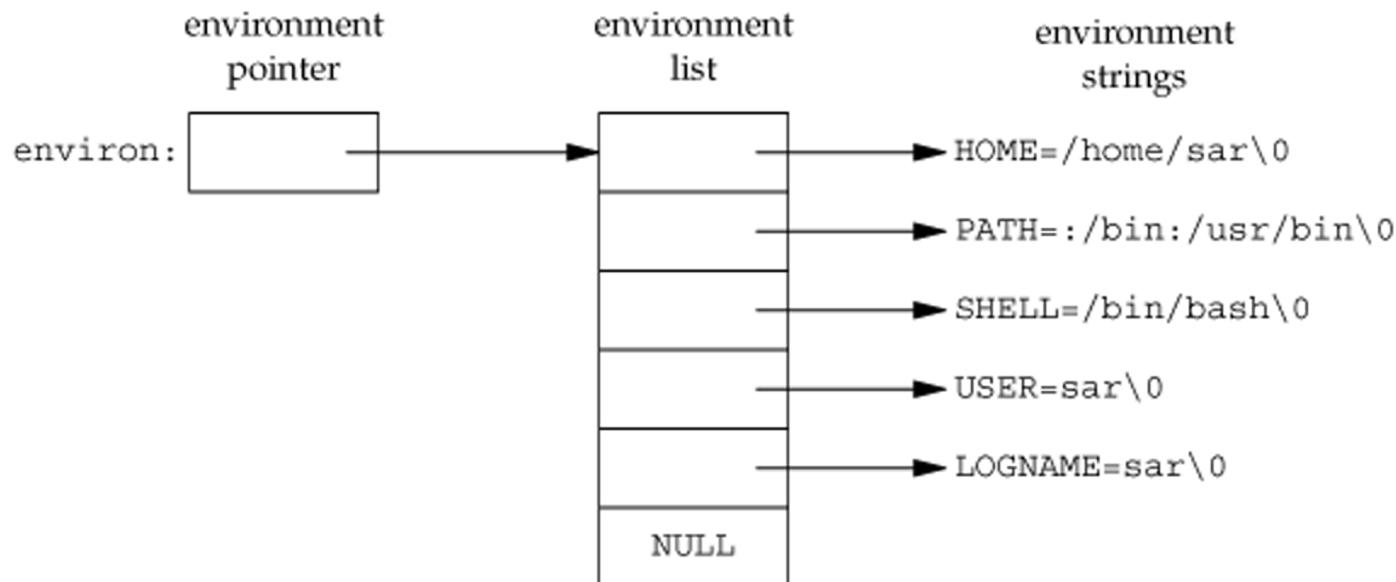
List of environment variable functions

Function	ISO C	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10
getenv	•	•	•	•	•	•
putenv		XSI	•	•	•	•
setenv		•	•	•	•	
unsetenv		•	•	•	•	
clearenv				•		

Environment List

Access environment variables directly

- `int main(int argc, char *argv[], char *envp[]);`
- `extern char **environ;`



Environment Functions

Prototypes of functions to manipulate environment variables

```
#include <stdlib.h>

char *getenv(const char *name);
int putenv(char *string);

int setenv(const char *name, const char *value, int overwrite);
int unsetenv(const char *name);
int clearenv(void);
```

Environment List Operations

Delete an entry

- This is simple, just free a string and move all subsequent pointers down one

Modify an entry

- If new-size \geq old-size, just overwrite the old one
- If new-size $>$ old-size, allocate a new space the new variable and make the pointer point to the new location

Add an entry

- Add for the 1st time, allocate a new space for the entire list
- Add for non-1st time, reallocate a larger space for the entire list

Common Environment Variables (1/3)

Variable	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
COLUMNS	•	•	•	•	•	Terminal width
DATEMASK	XSI		•	•	•	getdate(3) template file pathname
HOME	•	•	•	•	•	Home directory
LANG	•	•	•	•	•	Name of locale
LC_ALL	•	•	•	•	•	Name of locale
LC_COLLATE	•	•	•	•	•	Name of locale for collation
LC_CTYPE	•	•	•	•	•	Name of locale for character classification
LC_MESSAGES	•	•	•	•	•	Name of locale for messages

Common Environment Variables (2/3)

Variable	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
LC_MONETARY	•	•	•	•	•	Name of locale for monetary editing
LC_NUMERIC	•	•	•	•	•	Name of locale for numeric editing
LC_TIME	•	•	•	•	•	Name of locale for date/time formatting
LINES	•	•	•	•	•	Terminal height
LOGNAME	•	•	•	•	•	Login name
MSGVERB	XSI	•	•	•	•	fmtmsg(3) message components to process
NLSPATH	•	•	•	•	•	Sequence of templates for message catalogs

Common Environment Variables (3/3)

Variable	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
PATH	•	•	•	•	•	List of path prefixes to search for executable file
PWD	•	•	•	•	•	Absolute pathname of current working directory
SHELL	•	•	•	•	•	Name of user's preferred shell
TERM	•	•	•	•	•	Terminal type
TMPDIR	•	•	•	•	•	Pathname of directory for creating temporary files
TZ	•	•	•	•	•	Time zone information

Memory Layout of a Program

Text segment

- Machine instructions

Initialized data segment

- `int maxcount = 100;`

Uninitialized data segment (bss)

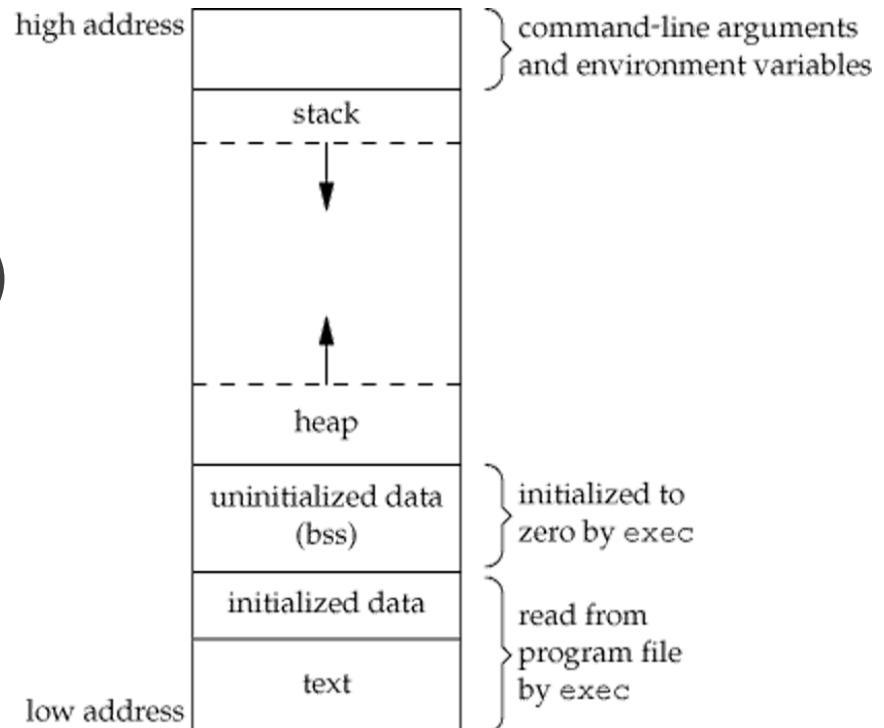
- `long sum[1000];`

Stack

- Local variables, function call states

Heap

- Dynamic allocated memory



Guess the Number

```
chuang@mjo:~/unix_prog/procenv — 57x24
7 int main() {
8     struct {
9         char buf[16];
10        int answer;
11    } d;
12
13    setvbuf(stdout, NULL, _IONBF, 0);
14    srand(time(0) ^ getpid());
15
16    d.answer = rand() % 10000;
17
18    printf("Guess the number: ");
19    if(fgets(d.buf, 20, stdin) != NULL) {
20        int g = strtol(d.buf, NULL, 0);
21        printf("Your guess is %d\n", g);
22        if(g == d.answer) {
23            printf("Bingo!\n");
24        } else {
25            printf("No no no ... \n");
26        }
27    }
28    return 0;
29 }
```

29 , 2

85%

Read Sizes of an Executable Binary

The size (1) command

```
$ size /usr/bin/gcc /bin/sh
  text      data      bss      dec      hex filename
203913      2152     2248   208313   32db9  /usr/bin/gcc
704028     19268    19736   743032   b5678  /bin/sh
```

Shared Libraries

Most UNIX systems today support shared libraries

Shared libraries remove the common library routines from the executable file

Maintain a single copy of the library routine somewhere in memory that all processes reference

- Reduce the size and memory requirement of each executable file
- But It may add some runtime overhead

Another advantage of shared libraries

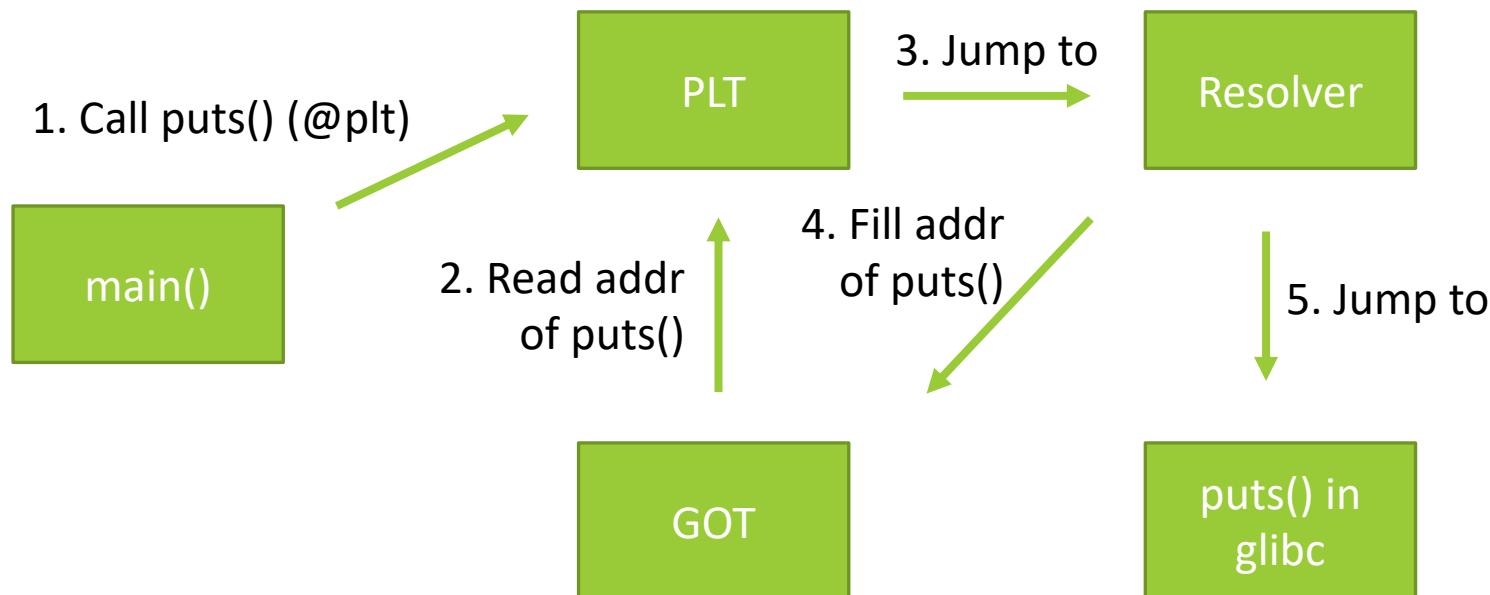
- Library functions can be replaced with new versions without having to relink every program that uses the library
- But it might also be a security flaw

Resolve Functions in Shared Libraries: call puts() – 1st Time

GOT – global offset table: Store the "real address" of a function

PLT – procedure linkage table: Call the real address, or resolve it!

Resolver – Procedures to resolve the "real address" of a function

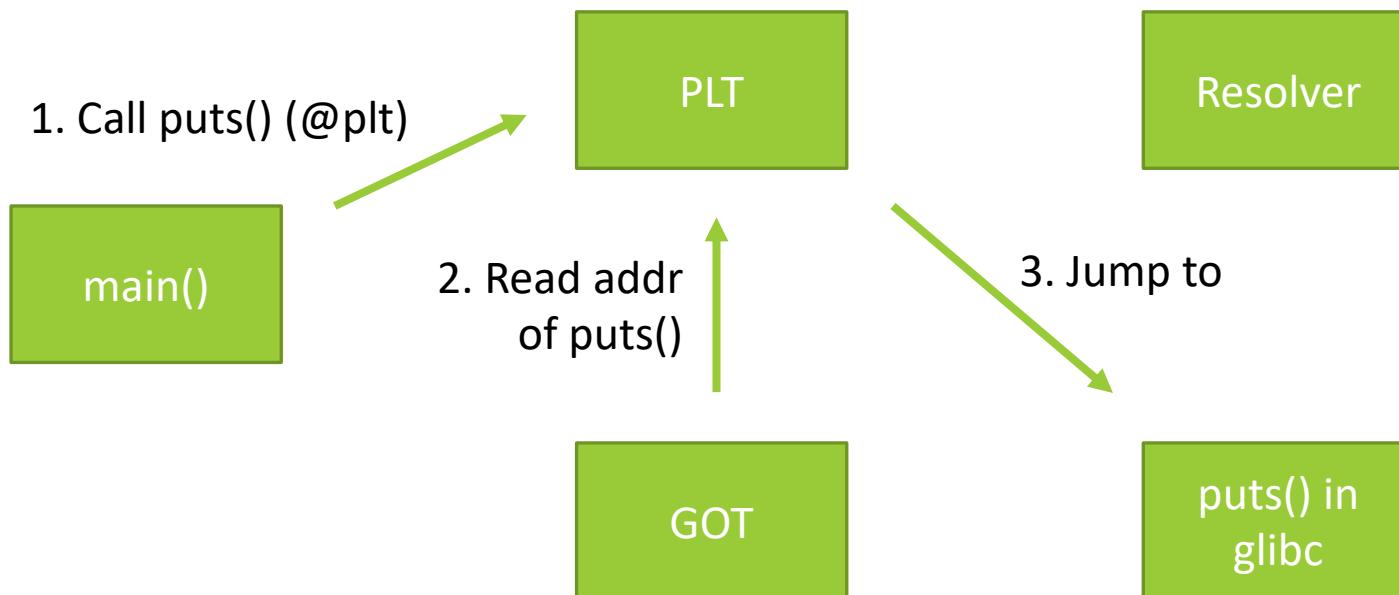


Resolve Functions in Shared Libraries: call puts() – 2nd Time +

GOT – global offset table: Store the "real address" of a function

PLT – procedure linkage table: Call the real address, or resolve it!

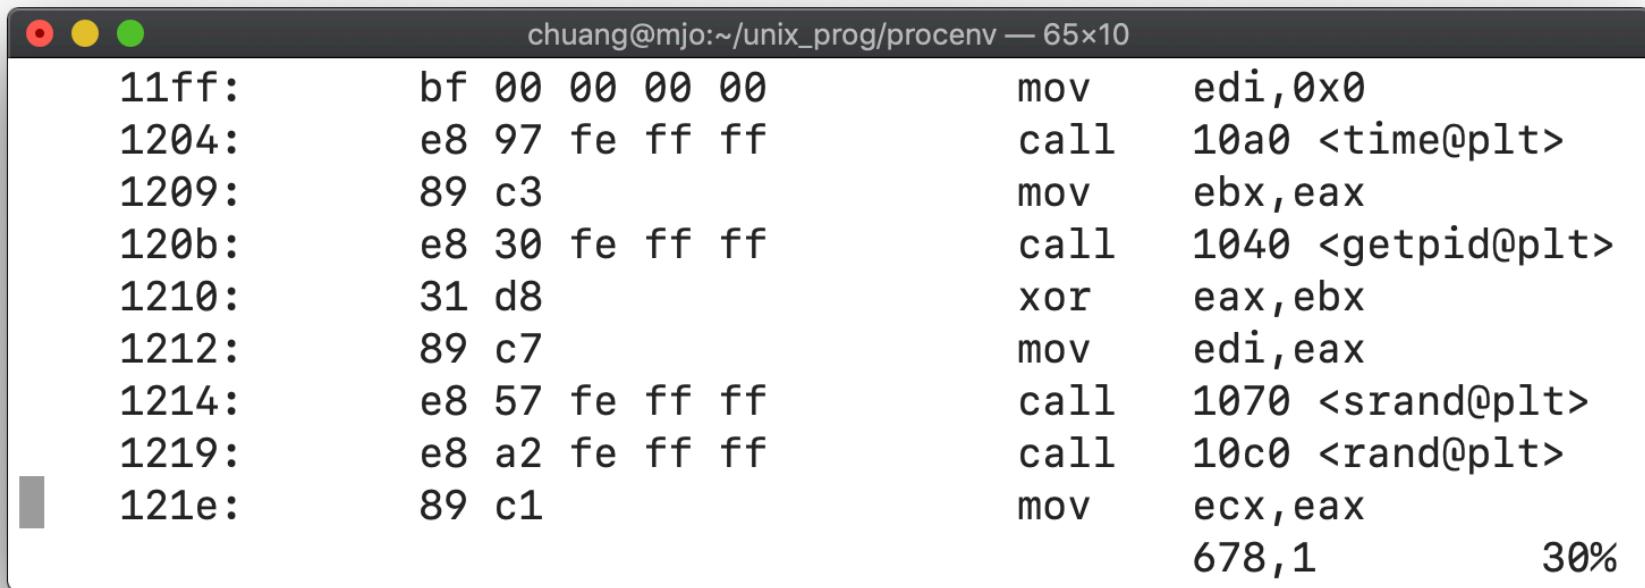
Resolver – Procedures to resolve the "real address" of a function



Resolving Functions in Shared Libraries – Sample: call time()

14

```
srand(time(0) ^ getpid());
```



The screenshot shows a terminal window with the title "chuang@mjo:~/unix_prog/procenv — 65x10". The window displays assembly code with its corresponding opcodes. The assembly code is:

```
11ff: bf 00 00 00 00          mov    edi,0x0
1204: e8 97 fe ff ff          call   10a0 <time@plt>
1209: 89 c3                   mov    ebx,eax
120b: e8 30 fe ff ff          call   1040 <getpid@plt>
1210: 31 d8                   xor    eax,ebx
1212: 89 c7                   mov    edi,eax
1214: e8 57 fe ff ff          call   1070 <srand@plt>
1219: e8 a2 fe ff ff          call   10c0 <rand@plt>
121e: 89 c1                   mov    ecx,eax
```

At the bottom right of the terminal window, there is a progress bar indicating "678,1" and "30%".

Resolving Functions in Shared Libraries – Sample: GOT & PLT

```
chuang@mjo:~/unix_prog/procenv — 103x11
00000000000010a0 <time@plt>:
 10a0: ff 25 aa 2f 00 00      jmp    QWORD PTR [rip+0x2faa]      # 4050 <time@GLIBC_2.2.5>
 10a6: 68 07 00 00 00      push   0x7
 10ab: e9 70 ff ff ff      jmp    1020 <.plt>

00000000000010c0 <rand@plt>:
 10c0: ff 25 9a 2f 00 00      jmp    QWORD PTR [rip+0x2f9a]      # 4060 <rand@GLIBC_2.2.5>
 10c6: 68 09 00 00 00      push   0x9
 10cb: e9 50 ff ff ff      jmp    1020 <.plt>

PLT
574,0-1 25%
```

```
chuang@mjo:~/unix_prog/procenv — 103x13
4049: 10 00      adc    BYTE PTR [rax],al
404b: 00 00      add    BYTE PTR [rax],al
404d: 00 00      add    BYTE PTR [rax],al
404f: 00 a6 10 00 00 00      add    BYTE PTR [rsi+0x10],ah
4055: 00 00      add    BYTE PTR [rax],al
4057: 00 b6 10 00 00 00      add    BYTE PTR [rsi+0x10],dh
405d: 00 00      add    BYTE PTR [rax],al
405f: 00 c6      add    BYTE PTR [rax],al
4061: 10 00      adc    BYTE PTR [rax],al
4063: 00 00      add    BYTE PTR [rax],al
4065: 00 00      dd     BYTE PTR [rax],al

Ox10a6
Ox10c6

GOT
1174,1-8 52%
```

Resolving Functions in Shared Libraries – Sample: DEMO

Relevant GDB commands

Display information

- x/30i – dump assembly instructions, e.g., x/30i main
- x/gx – dump a single 64-bit value in hexadecimal format
- tui enable – enable gdb's text user interface, "Ctrl-x o" to switch focus
- layout asm – change TUI layout to display assembly codes

Control flow

- b – set break points, e.g, b *main+59
- run – run the program
- start – run the first line of the program, and then paused
- starti – run the first instruction of the program, and then paused
- si – step one instruction
- ni – step one instruction, but proceed through subroutine calls

Compile Static and Dynamic Program

A simple program that just print “Hello, World!”

```
$ gcc h1.c -o h1
$ gcc h2.c -o h2 -static
$ ls -la h1 h2
-rwxrwxr-x 1 chuang chuang 9564 Mar 13 11:48 h1
-rwxrwxr-x 1 chuang chuang 878192 Mar 13 11:48 h2
$ size h1 h2
      text      data       bss      dec      hex filename
    896        264         8     1168      490   h1
 499650      1928     6948   508526    7c26e   h2
```

Library Injection

Functions referenced to shared libraries can be overridden

- The LD_PRELOAD environment variable
- Usage:

```
LD_PRELOAD=/path/to/the/injected-shared-object {program}
```

Library injection does not work for `suid`/`sgid` executables

Library Injection Example

Suppose we are going to hijack the getuid() function

- This is commonly used in tools like fake-root

The original program (getuid.c)

```
int main() {
    printf("UID = %d\n", getuid());
    return 0;
}
```

The injected library (inject1.c)

```
#include <stdio.h>
#include <sys/types.h>

uid_t getuid(void) {
    fprintf(stderr, "injected getuid, always return 0\n");
    return 0;
}
```

Library Injection Example (Cont'd)

Compile the programs and the libraries

```
$ gcc -o getuid -Wall -g getuid.c
$ gcc -o inject1.so -shared -fPIC inject1.c -ldl
```

- The first command produces the getuid program
- The second commands generates the inject1.so (shared) library

Run the example

```
$ ./getuid                      # no injection
UID = 1000
$ LD_PRELOAD=./inject1.so ./getuid # injected
injected getuid, always return 0
UID = 0
```

More on Library Injection

But we still want the original function to work properly

We have to locate the original function

```
#include <dlfcn.h>

void *dlopen(const char *filename, int flag);
char *dlerror(void);
void *dlsym(void *handle, const char *symbol);
int dlclose(void *handle);
```

You may have to link with -ldl option

Revised Library Injection Example

We would like to know the real UID internally (inject2.c)

```
#include <dlfcn.h>
#include <stdio.h>
#include <sys/types.h>

static uid_t (*old_getuid)(void) = NULL; /* function pointer */

uid_t getuid(void) {
    if(old_getuid == NULL) {
        void *handle = dlopen("libc.so.6", RTLD_LAZY);
        if(handle != NULL)
            old_getuid = dlsym(handle, "getuid");
    }
    fprintf(stderr, "injected getuid, always return 0\n");
    if(old_getuid != NULL)
        fprintf(stderr, "real uid = %d\n", old_getuid());
    return 0;
}
```

Revised Library Injection Example (Cont'd)

Compile the programs and the libraries (again)

```
$ gcc -o getuid -Wall -g getuid.c  
$ gcc -o inject2.so -shared -fPIC inject2.c -ldl
```

- The first command produces the getuid program
- The second commands generates the inject2.so (shared) library

Run the example

```
$ ./getuid                      # no injection  
UID = 1000  
$ LD_PRELOAD=./inject2.so ./getuid # injected  
injected getuid, always return 0  
real uid = 1000  
UID = 0
```

Determine Library Injection Possibility

No SUID/SGID enabled

Not a statically linked binary

Examples of the dynamic/static linked hello-world example

- The file command

```
$ file h1 h2
h1: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked (uses shared
libs), for GNU/Linux 2.6.24, BuildID[sha1]=e32f08cfbdda94d57273829c2bfd535d8fbe626d, not
stripped
h2: ELF 64-bit LSB executable, x86-64, version 1 (GNU/Linux), statically linked, for
GNU/Linux 2.6.24, BuildID[sha1]=2748d80822e76d183d0ef5633c0b784527727c7a, not stripped
```

- The ldd command

```
$ ldd h1 h2
h1:
    linux-vdso.so.1 => (0x00007ffe7d3d5000)
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f1bc2150000)
    /lib64/ld-linux-x86-64.so.2 (0x00007f1bc2515000)
h2:
    not a dynamic executable
```

Determine Library Injection Possibility (Cont'd)

Use symbols from a shared library

The `nm` command

Example: static VS dynamic linked symbols

```
$ gcc -o getuid -Wall -g getuid.c           # dynamically linked
$ gcc -o getuid_s -Wall -g getuid.c -static  # statically linked
$ nm getuid | grep getuid
                           U getuid@@GLIBC_2.2.5      # getuid is unknown
$ nm getuid_s | grep getuid
000000000433590 W getuid                   # getuid is known (but weak)
000000000433590 T __getuid                 # the getuid implementation
```

Symbols can be stripped using the `strip` command

Memory Allocation

ISO C memory allocation functions

`void *malloc(size_t size);`

- Allocates a specified number of bytes of memory
- The initial value of the memory is indeterminate

`void *calloc(size_t nobj, size_t size);`

- Allocates space for a specified number of objects of a specified size
- The space is initialized to all 0 bits

`void *realloc(void *ptr, size_t newsize);`

- Increases or decreases the size of a previously allocated area
- It may involve moving the previously allocated area somewhere else, to provide the additional room at the end
- The initial value of increased memory is indeterminate

Memory Allocation (Cont'd)

Allocated memory can be released by free()

The allocation routines are usually implemented with the sbrk(2) system call

This system call expands (or contracts) the heap of the process

- However, most versions of malloc and free never decrease their memory size
- The space that we free is available for a later allocation
- The freed space is usually kept in the malloc pool, not returned to the kernel

The alloca Function

A special memory allocation function – alloca

```
#include <alloca.h>
void *alloca(size_t size);
```

alloca() allocate memories in **stack frames** of the current function call

So you don't have to free() the memory – it is released automatically after the execution of the current function returns

May be not supported by your system, but modern UNIXes supports the function (Linux, FreeBSD, Mac OS X, Solaris)

Pros: might be faster (than malloc), no need to free, easier to work with setjmp/longjmp

Cons: Portability

setjmp and longjmp Function

The reserved keyword "goto" can be used only in the same function

We cannot goto a label that is in another function

Instead, we must use the setjmp and longjmp functions to perform this type of branching

Typical Program Skeleton for Command Processing

```
#include "apue.h"
#define TOK_ADD      5
void    do_line(char *);
void    cmd_add(void);
int     get_token(void);

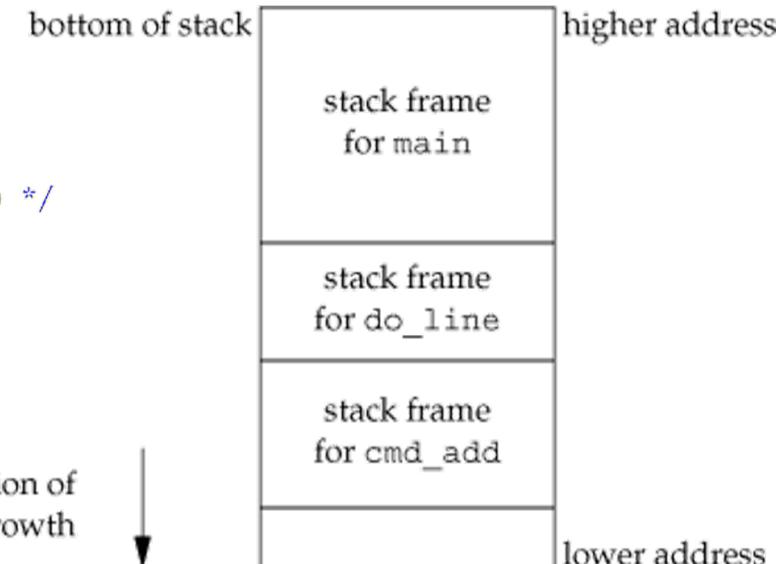
int main(void) {
    char    line[MAXLINE];
    while (fgets(line, MAXLINE, stdin) != NULL)
        do_line(line);
    exit(0);
}

char    *tok_ptr;          /* global pointer for get_token() */
void    do_line(char *ptr) { /* process one line of input */
    int     cmd;
    tok_ptr = ptr;
    while ((cmd = get_token()) > 0) {
        switch (cmd) { /* one case for each command */
            case TOK_ADD: cmd_add(); break;
        }
    }
}

void    cmd_add(void) {
    int     token;
    token = get_token(); /* rest of processing for this command */
}

int    get_token(void) {
    /* fetch next token from line pointed to by tok_ptr */
}
```

What if we encounter an error in cmd_add and would like to jump back to the main function for processing the next line?



The Solution for Jumping Across Functions

Set the jump back position

- `int setjmp(jmp_buf env);`
- `env` is usually a global variable – has to be accessed from both the `setjmp` side and the `longjmp` side
- Returns: 0 if called directly, or nonzero if returning from a call to `longjmp`

Jump back

- `void longjmp(jmp_buf env, int val);`
 - The '`val`' will be returned from `setjmp`
 - If `val` is 0, it will be replaced by 1

Using setjmp and longjmp

```
#include "apue.h"
#include <setjmp.h>

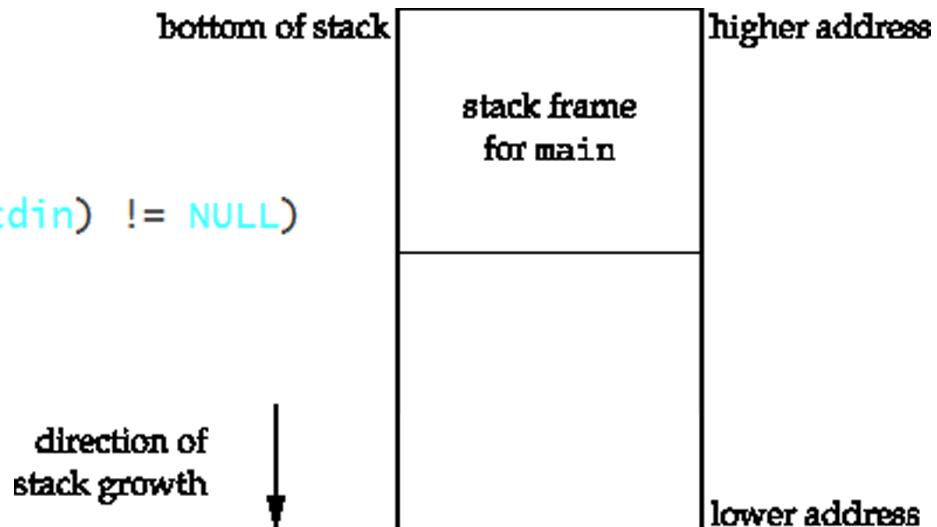
#define TOK_ADD      5

jmp_buf jmpbuffer;

int main(void) {
    char line[MAXLINE];
    if (setjmp(jmpbuffer) != 0)
        printf("error");
    while (fgets(line, MAXLINE, stdin) != NULL)
        do_line(line);
    exit(0);
}

void cmd_add(void) {
    int token;
    token = get_token();
    if (token < 0)          /* an error has occurred */
        longjmp(jmpbuffer, 1);
    /* rest of processing for this command */
}
```

Stack after jumped back



Restoration of Variables (1/4)

Type of variables

- Automatic, e.g., `[auto] int autoVal;`, the default
- Register, e.g., `register int regVal;`, store in register if possible
- Volatile, e.g., `volatile int volVal;`, store in memory

What are the values of variables after jumped back?

- It depends
- Most implementations do not try to roll back these automatic variables and register variables
- The standards say only that their values are indeterminate
- If you have an automatic variable that you do not want to be rolled back, define it with the volatile attribute
- Variables that are declared global or static are left alone when longjmp is executed
- **In short: variables in register – restored; variables in memory – kept**

Restoration of Variables (2/4)

```
#include "apue.h"
#include <setjmp.h>
static void      f1(int, int, int, int);
static void      f2(void);
static jmp_buf   jmpbuffer;
static int       globval;

int main(void) {
    int          autoval;
    register int  regival;
    volatile int  volaval;
    static int    statval;
    globval = 1; autoval = 2; regival = 3; volaval = 4; statval = 5;
    if (setjmp(jmpbuffer) != 0) {
        printf("after longjmp:\n");
        printf("globval = %d, autoval = %d, regival = %d,"
               " volaval = %d, statval = %d\n",
               globval, autoval, regival, volaval, statval);
        exit(0);
    }
    // change variables after setjmp, but before longjmp.
    globval = 95; autoval = 96; regival = 97; volaval = 98; statval = 99;
    f1(autoval, regival, volaval, statval); /* never returns */
    exit(0);
}

static void f1(int i, int j, int k, int l) {
    printf("in f1():\n");
    printf("globval = %d, autoval = %d, regival = %d,"
           " volaval = %d, statval = %d\n", globval, i, j, k, l);
    f2();
}

static void f2(void) { longjmp(jmpbuffer, 1); }
```

Restoration of Variables (3/4)

Rules for variable restoration

- Variables stored in memory will have values as of the time of calling `longjmp`
- Variables in the CPU and floating-point registers are restored to their values when `setjmp` was called

Hence,

- auto variables may be indeterminate, it depends on compiler implementations
- register variables are restored to the value of “before calling `setjmp`”
- volatile variable are restored to the value of “before calling `longjmp`”

Restoration of Variables (4/4)

Set 1,2,3,4,5 → setjmp → Set 95,96,97,98,99 → longjmp → ?

- No optimization: gcc places everything in memory
- Full optimization: auto/register variables are placed in registers

```
$ gcc fig7.13-testjmp.c -I../include -o t1      compile without any optimization
$ gcc fig7.13-testjmp.c -I../include -o t2 -O    compile with full optimization
$ ./t1
in f1():
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
after longjmp:
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
$ ./t2
in f1():
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
after longjmp:
globval = 95, autoval = 2, regival = 3, volaval = 98, statval = 99
```

Process Resource Limits

Every process has a set of resource limits

Resource limits are usually initialized by a parent process and inherited by its child processes

The getrlimit and setrlimit functions

```
#include <sys/time.h>
#include <sys/resource.h>
int getrlimit(int resource, struct rlimit *rlim);
int setrlimit(int resource, const struct rlimit *rlim);
```

The rlimit structure

```
struct rlimit {
    rlim_t rlim_cur; /* Soft limit */
    rlim_t rlim_max; /* Hard limit (ceiling for rlim_cur) */
};
```

Partial List of Process Resources

Limit	XSI	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10
RLIMIT_AS	•	•	•		•
RLIMIT_CORE	•	•	•	•	•
RLIMIT_CPU	•	•	•	•	•
RLIMIT_DATA	•	•	•	•	•
RLIMIT_FSIZE	•	•	•	•	•
RLIMIT_MEMLOCK		•	•	•	
RLIMIT_NOFILE	•	•	•	•	•
RLIMIT_NPROC		•	•	•	
RLIMIT_RSS		•	•	•	
RLIMIT_SBSIZE		•			
RLIMIT_STACK	•	•	•	•	•
RLIMIT_VMEM					•

Example to Dump Resource Limits

See code fig7.16-getrlimit.c

```
$ ./fig7.16-getrlimit
RLIMIT_AS          (infinite)  (infinite)
RLIMIT_CORE        1024000000  (infinite)
RLIMIT_CPU         (infinite)  (infinite)
RLIMIT_DATA        (infinite)  (infinite)
RLIMIT_FSIZE       (infinite)  (infinite)
RLIMIT_LOCKS       (infinite)  (infinite)
RLIMIT_MEMLOCK     65536       65536
RLIMIT_NOFILE      1024        4096
RLIMIT_NPROC       96120       96120
RLIMIT_RSS         (infinite)  (infinite)
RLIMIT_STACK       8388608   (infinite)
```

Example to Dump Resource Limits

Limits	Description
RLIMIT_CORE	The maximum size in bytes of a core file. A limit of 0 prevents the creation of a core file.
RLIMIT_MEMLOCK	The maximum amount of memory in bytes that a process can lock into memory using mlock(2).
RLIMIT_NOFILE	The maximum number of open files per process.
RLIMIT_NPROC	The maximum number of child processes per real user ID.
RLIMIT_STACK	The maximum size in bytes of the stack.

Q & A
